MelTS

Melody Translation System

Nicole Limtiaco
limni@seas.upenn.edu
Univ. of Pennsylvania
Philadelphia, PA

Rigel Swavely
rigel@seas.upenn.edu
Univ. of Pennsylvania
Philadelphia, PA

ABSTRACT

MelTS is an automatic harmonization system that creates multi-part arrangements in the style of the data on which it was trained. The system approaches the problem of harmonization from a machine tranlsation perspective, modeling the melody of a song as the source language and each harmony as a target language. The approach stands in contrast to previous approaches to the harmonization problem, which have primarily taken two forms: satisfying rules provided by music theory or predicting the chord under the melody. A major benefit of the MelTS approach is that the style of the harmony voices can be learned directly from the training data, just as vocabulary and structure can be learned from a parallel corpus of natural languages. In particular, generating harmony lines probabilistically and individually, as opposed to by rule satisfaction or as a consequence of a predicted chord, allows for the production of more natural arrangements without the need for simplifying assumptions about the structure of the music to be produced.

1. INTRODUCTION

Multi-part musical arrangements are a cornerstone of many musical styles. From choirs and string quartets to barbershop and a-cappella groups, musicians are constantly producing creative arrangements to harmonize with their favorite melodies. The automatic harmonization problem is one in which machines are put to this uniquely human task of composing music to support a melody.

Formally, a *melody* is a sequence of input *notes*, where each note contains information about pitch and timing. A *harmony* is a sequence of output notes which are produced with the constraint that the sequence of notes supports the input melody. Both *melodies* and *harmonies* may be referred to as *parts* in an arrangement. The term *voice* is defined as some category of *parts* with a certain set of identifying characteristics. For example, the bass lines in rock music are one type of voice and the soprano solos in an opera are another. The aim of the automatic harmonization problem is to, given a melody part, generate parts in n different harmony voices that, when played together, sound coherent and pleasant.

MelTS, a melody translation system, is a proof of concept endeavor aimed at reducing the automatic harmonization problem to a machine translation problem, viewing the melody voice as the source language and the harmony voices as the targets. Such an approach allows the style of the desired harmonies to be learned from a set of musical training data, just as the intricacies of a natural language can be learned from a parallel corporus.

2. RELATED WORK

2.1 Automatic Harmonization

Automatic harmonization is a subset of the automatic musical composition problem, which dates as far back as the field of artificial intelligence itself. Perhaps the earliest work in automatic composition is Hiller and Isaacson's *Illiac Suite* [5], which is widely accepted as being the first musical piece composed by an electronic computer. Hiller and Isaacson used a generate-and-test method that generated musical phrases psuedo-randomly and kept only those that adhered to a set of music-theory-inspired heuristics.

Staying in line with the use of musical rules, Ebcioğlu [3] provided a break-through in the specific field of automatic harmonization through his CHORAL system. CHORAL, bolstered by about 350 music-theory and style rules expressed in first order logic, performed the task of writing four-part harmonies in the style of J.S. Bach. With these logical predicates, Ebcioğlu reduced the problem of composing a harmony to a simple constraint satisfaction problem. Similar later works, notably by Tsang & Aitkin [12], also crafted the harmonization problem as a problem in constraint satisfaction, but with a significantly smaller amount ($\sim$20) of musical rules. The result of these constraint-based works were musically sensible; however, crafting the constraints such that the output is musical requires deep, human knowledge about music in general and about the style of music to be produced in particular.

More recent works have put data-driven methods into use in order to infer the patterns that govern real compositions. A simple case-based model implemented by Sabater *et al.* [9] was built to generate accompanying chords for a melody line. To a choose a harmony chord for a given context of previous harmony chords and the currently sounding melody note, the system would check a case base to see if any cases had the same harmony context and melody note and use the corresponding harmony chord if such a match were found. Musical heuristics were used to generate a chord if no match was found in the case base. An automatic harmonizer utilizing neural networks was also built by Hild *et al.* [4] to produce a harmony chord for a given melody quarter beat. Input features to the network included harmony context, current melody pitch, and whether or not the beat was stressed.

As these examples show, the previous harmony context and the melody pitch are important signals in deciding what the current harmony phrase should be. Many works have

been conducted that model these signals using n-gram Markov Models. A Markov model assigns probabilities to some event C_t conditioned on a limited history of previous events $[C_{t-1} \ldots C_{t-n}]$, implictly making the assumption that the event C_t depends only on a small amount of information about its context. For example, a system called MySong produced by Simon *et al.* [11] generates chord accompaniment given a vocalized melody by using a 2-gram Markov model for the harmony context and a 2-gram Markov model for the melody context. A similar system implented by Scholz *et al.* [10], which also generates chord accompanimants, experimented with 3-gram to 5-gram Markov models and incorporated smoothing techniques commonly seen in NLP to account for cases in which the model has not seen a context that is present in the test data. Most recently, Raczyński *et al.* [8] uses discriminative Markov models that model harmony context, melody context, and additionaly the harmony relationship to the tonality.

A recent senior design project out of the University of Pennsylvania by Cerny *et al.* [1] also used melody pitch and previous harmony context as their main signals for determining the next harmony chord to generate. However, they used these signals as inputs to an SVM classifier, as opposed to training a Markov model.

2.2 Machine Translation

MelTS reduces the automatic harmonization problem to a machine translation problem. However, automatic harmonization differs in an interesting way from standard machine translation applications in that it requires translation from one language (melody voice) to many languages (n harmony voices) instead of one language to one other language. Looked at from a different perspective, the problem can be seen as a sequence of translations from many source languages to one target language. The sources are the melody voice and all the harmony voices, if any, that have been produced so far; the target is the harmony voice we would like to produce next. Though limited work has been done on the multi-source translation problem, Och and Ney [7] produced a work describing several ways of altering standard methods to deal with multiple sources.

3. SYSTEM MODEL

3.1 Motivating the Machine Translation Approach

The previous data driven approaches applied to automatic harmonization all share in the fact that they predict the chords to be played under the input melody. Automatic harmonization via chord prediction imposes two major limitiations: a limited set of produceable chords and a lack of fluid movement within the individual parts.

In the best case, chord prediction will alow only the generation of those chords seen in the training data. In more restrictive set-ups where classification algorithms are used, the chord predicted is based on a classifier choosing from a relatively small number of chord classes. Some chord prediction systems do not generate individual parts at all, but rather view the harmony output as just the underlying chord sequence generated. If individual parts are produced, the notes for each part are chosen from the predicted chord, not based on the context of previous notes in that part. Chord prediction further limits interesting musical structure since

some assumptions must be made about when the predicted chord is to be played. This is a far cry from actual musical arrangements, whose parts can move independently of each other and at times produce non-conventional chords. In fact, this approach differs greatly from how actual composers create multi-part arrangements, creating musical phrases in each part rather than purely choosing chords to sound on each beat.

Predicting the note sequences for individual parts in such a way that will allow them to sound consont when played together can bypass many of the pitfalls of chord prediction. Such an approach will allow unseen chords to be produced because there is no restriction on which groups of notes can sound simultaenously, only on which notes can be produced. Furthermore, the system can be made to encourage interesting movement within the individual parts, for example by taking the part's context into account and by allowing rhythmic variation. Machine translation techniques are used to achieve these goals of simultaneous consonance and part fluidity.

There are two main analogies between natural languages and musical voices that motivate the usage of machine translation techniques for automatic harmonization. The first analogy is that both natural languages and musical voices have a sense of "word translations". Just as there may be several words in the target language that could be sensible translations of a given word in the source langauge, there are also several notes in the harmony voice that can harmonize well with an input melody note. Importantly, however, it is not the case that any note can harmonize well with the melody note. Some notes will sound dissonant when played together and still other notes may not be in the harmony voice at all, since harmony voices may have specific note ranges. The second analogy is that both natural languages and musical voices have a concept of "fluency", the idea that only certain strings of tokens (i.e. words or notes) are sensible. For example, the statement "colorless green ideas sleep furiously" contains all English words but is unlikely to be understood by an English speaker since that string of words is meaningless based on the rules of the language. Similarly, a random sequence of notes may not sound sensible in the context of its harmony voice, if the notes are recognized as music at all. Since machine translation techniques can produce novel natural language utterances based on the properties of word translations and fluency, it follows that novel musical parts can be generated with the same techniques.

3.2 System Design Overview

In line with the analogies explained above, two probabilistic models commonly seen in statistical machine translation are employed: the language model and the translation model. The language model, denoted as $L(H)$, provides a probability for a sequence of pitches in a harmony part. The translation model is made up of two sub-models: the note model, denoted $N(M|H)$, and the phrase model, denoted $Ph(M|H)$. The note model provides a probability of harmonization between melody and harmony pitches. The phrase model provides a probability of harmonization between two-note musical phrases, consisting of pitch and rhythm information, between the meldoy and the harmony. Specifically, the models are given by the formulas below:

$$L(H) = \Pi_{i=1}^{l} P[pitch(h_i)|pitch(h_{i-1}) \ldots pitch(h_{i-n})]$$
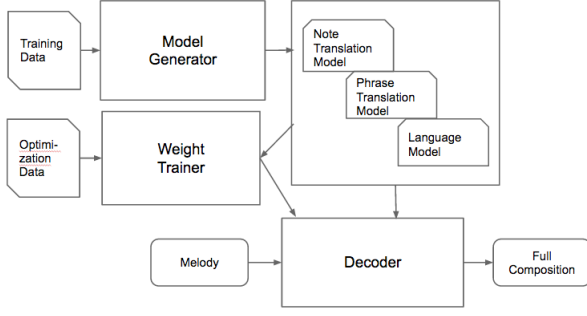
Figure 1: Overview of System Design

$$N(M|H) = \Pi_{i=1}^{l} P[pitch(m_i)|pitch(h_j)]$$

$$Ph(M|H) = \Pi_{i=1}^{l-1} P[(m_i, m_{i+1})|(h_j \ldots h_{j+k})]$$

The model scores for a generated harmony are combined into a total score using a weighted log-linear model. In general, a log-linear model computes a sum of the features' values each multiplied by its corresponding weight. The model adapted for these purposes is defined below:

$$S(H|M) = [w_L \cdot log(L(H))]+$$

$$[w_{Ph} \cdot log(Ph(M|H))] + [w_N \cdot log(N(M|H))]$$

A learning algorithm is used to determine the values of $L(H)$, $N(M|H)$, and $Ph(M|H)$ based on a set of music training data. The weights corresponding to the models are determined with a weight-optimization algorithm that utilizes a separate set of music training data. The details of the alrogithms are discussed in the implementation section.

Given a melody part M, the goal is to find some harmony part H that maximizes the score $S(H|M)$. A decoder utilizing the language and tranlsation models is employed to find the harmony part in the search space that maximizes the weighted log-linear score for the given melody part. Figure 1 shows graphical overview of the model just described.

Up until this point, it has been assumed that the translation is into only one harmony voice. However, the goal is to produce a comple n-part arrangement. In order to do this, the decoding portion of the system is viewed as a sequence of multi-source translations where the source languages are the melody voice and all of the harmony voices generated so far. The generated harmony maximizes a weighted log-linear score that incorporates its language model score and the scores given by each translation model from one of the existing voices to it. Specifically, given a set V of source voices, the multi-source version of problem finds a harmony part H such that

$$H = \underset{H}{argmax}\ S(H|V) = [w_L \cdot log(L(H))]+$$

$$[w_N \cdot \Sigma_{v \in V} log(N(v|H))] + [w_{Ph} \cdot \Sigma_{v \in V} log(Ph(v|H))]$$

Viewing translation as a sequence of generation steps then introduces the problem of determining the best order for the parts to be generated. Since each part is constrained by the

parts generated previously, one would expect that the order in which the parts are generated will have an observable affect on the output. Indeed, our experiments have shown that the quality of the output varies greatly among different generation orderings. Details about how we choose an optimal ordering will be discussed in the implementation section.

4. SYSTEM IMPLEMENTATION

4.1 Data Collection and Interaction

To interact with music data, a third-party software called Music21[2] was used. Music21 is a library developed at MIT which provides a simple API for querying several types of music formats. MusicXML, a type of encoding for standard music notation, is the primary data format used. This format can be parsed by Music21 into Stream objects which can then be queried for any of the information contained in the notation, including key, tempo, and notes with pitch and timing. Additionally, the Music21 package has a library of over $2,000$ compositions, mostly written by classical composers. The package's extensive collection of Bach chorales — which each have soprano, alto, tenor, and bass voices — was used as the main source of training, optimization, and test data.

	Songs	S. Notes	A. Notes	T. Notes	B. Notes
Total	351	21062	25255	25691	25867
Major	185	11371	13518	13647	13846
Minor	166	9691	11737	12044	12021

Table 1: Number of songs and notes in the training data set. "S" stands for soprano, "A" for alto, "T" for tenor, and "B" for bass

4.2 Note Representation

Notes are represented in the models as strings of the form:

```
pitch -> {pitch class}{octave}
note -> {pitch}(,{pitch})*(:{quarter-length duration})?
```

Every model uses at least the pitch information for a note, so the pitch class and octave are necessary in the representation. For exmaple, the pitch representation of a middle C is `C4`. A chord, a group of notes that sound together in the same part, can be represented by repeating the pitch and octave portion of the representation. For example, a representation of a C-major chord would be `C4,E4,G4`. Since the phrase model incorporates timing information, the note representations in that model include the note's quarter-lengths appended to the pitches after a colon. For example, a middle C eighth note is represented as `C4:0.5`.

Rests, moments of silence in the composition, are considered as "notes" in our model. They are represented by the string `R`. Including rests makes it possible to harmonize a melody note with silence, thus offering some opportunity for rhythmic variation. However, including rests does pose the potential problem of producing too much silence, something which is generally not optimal for music generation. For example, imagine a language model of n-gram size 2 in which $P(R|R)$ and $P(R|R\ R)$ are both relatively high. Once one rest is produced, the system may very well produce rests

for the remainder of the song. Therefore, for fear of continually producing rests, all contiguous sequences of rest symbols are treated as one rest `R` in the model.

Measure bars are also modeled as "notes" in our system. They are represented by the string `BAR`. A special measure bar, the bar at the end of last measure in the song, is given a special representation: `END`. These notes are slightly different in that they are only present in our language model. `BAR` and `END` can only be harmonized with actual measure bars and song end in the given melody part, so there is no need to include them in the translation model. The motivation for including `BAR` and `END` in the language model is that they provide information about where in the composition the system is trying to produce notes. This information is helpful because some notes may be more likely to start or end measures and some sequences of notes may be more likely to end a song. Coming back to the analogies between languages and musical voices, the `BAR` and `END` symbols can be likened to punctuation marks, which can give very strong cues about which words to generate.

4.3 Model Generation

The models represent probabilities of events as nested dictionaries. The outer dictionary maps the given event, A, to the inner dictionary, which maps the unknown event, B, to the probability $P(B|A)$. In the case of the language model, A is the sequence of $(i-1)...(i-n)$ harmony pitches and B is the i^{th} harmony pitch, whereas in the note translation model, A and B are pitches in the harmony and melody, respectively, that sound simultaneously. Similarly in the phrase translation model, A and B are two-note sequences in the harmony and melody, respectively, that sound simultaneously.

The probabilities are calculated based on the counts of the events in the training compositions, plus smoothing techniques applied so that no event is assigned a zero probability. For now, the smoothing techniques are ommitted for ease of explanation. Formally, the language ngram probabilities are given by:

$$P(pitch(h_i)|pitch(h_{i-1})...pitch(h_{i-n})) =$$

$$\frac{count(pitch(h_i) \wedge pitch(h_{i-1}) \wedge ... \wedge pitch(h_{i-n}))}{count(pitch(h_{i-1}) \wedge ... \wedge pitch(h_{i-n}))}$$

The note translation probabilities are given by:

$$P(pitch(m_i)|pitch(h_j)) = \frac{count(pitch(m_i) \wedge pitch(h_j))}{count(pitch(h_j))}$$

and the phrase translation probabilities are given by:

$$P((m_i, m_{i+1})|(h_j \ldots h_{j+k})) = \frac{count((m_i, m_{i+1}) \wedge (h_j \wedge ... \wedge h_{j+k}))}{count((h_j \wedge ... \wedge h_{j+k}))}$$

The event $pitch(h_i) \wedge ... \wedge pitch(h_{i-n})$, as seen in the ngram probability formula, occurs whenever the pitches of notes $h_{i-n}, h_{i-n+1}, ..., h_i$ appear contiguously in that order in the haromny part.

The event $pitch(m_i) \wedge pitch(h_j)$, as seen in the note translation probability formula, occurs whenever the pitch of note m_i is sounding in the melody at the same time that the pitch of note h_j is sounding in the harmony, regardless of when either note begins or ends. To compute these counts, we iterate through each note m_i in the melody and update the counts for event $m_i \wedge h_j$ for all notes h_j that have any portion

Figure 2: Two phrase events are boxed. The original harmony part is displayed along with edited notes that line up with their corresponding melody phrase.

of their duration sounding at the same time as any portion of the duration of m_i.

The counts for the events $(m_i, m_{i+1}) \wedge (h_j \wedge \ldots \wedge h_{j+k})$, as seen in the phrase translation model probability formula, are gathered by moving a 2-note wide sliding window across the melody part and determing, for each 2-note melody phrase, the phrase of notes $(h_j \ldots h_{j+k})$ that are playing simultaneously in the harmony part. Since the phrase translation model takes timing information into account, the training algorithm insures that the melody and harmony phrases have the same duration. If the first note in the harmony phrase begins sounding before the first note in the melody phrase, its duration in the model is shortened so that it only sounds when the first melody phrase note sounds. Similarly, if the last note in the harmony phrase finishes after the last note in the melody phrase, its duration in the model is shortened so that if finishes at the same time as last note in the melody phrase. For example, take the melody phrase `[C4:1.0 E4:1.0]` and the harmony phrase `[E4:1.0 C4:0.5 C4:0.25 G4:1.0]` which starts 0.5 of a quarter-length before the melody phrase. The corresponding harmony phrase event will be `[E4:0.5 C4:0.5 C4:0.25 G4:0.75]`. See Figure 1 for a pictoral explanation.

LaPlace smoothing, also known as additive smoothing, is incorporated into the language model to avoid any events having zero probability. Specifically, for some smoothing constant α, the language ngram probabilities are given by:

$$P(pitch(h_i)|pitch(h_{i-1})...pitch(h_{i-n})) =$$

$$\frac{count(pitch(h_i) \wedge pitch(h_{i-1}) \wedge ... \wedge pitch(h_{i-n})) + \alpha}{count(pitch(h_{i-1}) \wedge ... \wedge pitch(h_{i-n})) + (48\alpha)}$$

The constant 48 in the equation above is meant to be the number of possible note values for the variable h_i. It is derived from the fact that there are 12 notes in each octave and that 4 is a very liberal estimate for how many octaves a particular voice can span.

The translation models use a simple smoothing constant, $\alpha = 1e - 10$, for when there is no probability associated with a translation event.

Lastly, the ngram size n and smoothing parameter α have been choosen as $n = 3$ and $\alpha = 1e - 10$. Qualitatively, it seems that these values are sufficient, though one could imagine optimizing these parameters to maximize the qual-

ity of the system output.

4.4 Weight Trainer

The weights for each model are tuned using the Powell algorithm [6] on a set of held out training data, separate from the data used to train the models. If the harmonies to be generated are for major songs, only the major compositions in the training data are used, and vice versa for minor songs. The algorithm tunes the weights to optimize an arbitrary metric; in this case, the weighted log-linear score was optimized. See Table 2 for information on the held out training set.

	Songs	S. Notes	A. Notes	T. Notes	B. Notes
Total	20	1178	1392	1413	1468
Major	11	658	779	793	822
Minor	9	520	613	620	646

Table 2: Number of songs and notes in the held out training data set used for weight tuning. "S" stands for soprano, "A" for alto, "T" for tenor, and "B" for bass

4.5 Decoder

The goal of the decoder is to find the harmony part H such that H maximizes the weighted log-linear score $S(H|M)$. However, the search space is much too large to enumerate all possible harmony parts in order to determine the best. Assuming the harmony voice spans no more than 4 octaves and that there are 12 notes in an octave, an exhaustive search of all harmony parts for a song with n melody notes would require the inspection of $(12 \cdot 4)^n = 48^n$ harmony parts. For an average length song comprised of 100 melody notes, the number of possible harmony parts comes out to $48^{100} \approx 1.3e + 168$.

To shrink the search space, harmony parts are generated four measures at a time. That is, for each four measure chunk, a set of k harmonies are generated. The cross-product of those k harmonies with the stored harmonies corresponding to the previous measures result in a new set of full-song harmony prefixes, out of which the top ranking j are stored for the next iteration.

The harmony generation process for a four-measure chunk is accomplished with a beam-seach decoder. Secifically, at every iteration t, the decoder will have a set of hypotheses S_t where each hypothesis is a sequence of notes with a duration equivalent to the duration of the melody part up to its $min(2t,l)^{th}$ note, where l is the total number of notes in the melody. The decoder is begun with the set S_0 containing only the empty hypothesis.

For each two-note phrase in the melody, the decoder finds all possible harmony phrases that can sound with that melody phrase, call it $m_i = (m_i^0, m_i^1)$. A portion of the harmony phrase possibilites comes from the corresponding harmony phrases for m_i in the phrase translation model. The other possibilities are two-note phrases which have the same rhythm as m_i, but whose pitches are determined by the note translation model. The cross product of the pitches corresponding to m_i^0 with the pitches corresponding to m_i^1 in the note translation model give the pitches for this portion of the possible harmony phrases. There are two important things to

note about generating possible harmonies from the phrase translation model. The first is that since the harmony and melody phrase pairs in the phrase translation model are always of the same duration, the harmony phrases can be added to the hypotheses without having to worry about duration mismatches between the hypotheses and the translated melody prefix. Secondly, the possible harmonies generated from the phrase translation model are the main source of rhythmic variation in the harmony generation algorithm. Those phrases are what prevents the output harmonies from being exactly lined up with the melody, leading to more natural-seeming arrangements.

Once the possible harmony phrases are retrieved, the decoder constructs S_{t+1} by examining each hypothesis in S_t. For each hypothesis $hyp \in S_t$, a new hypothesis is constructed for each possible harmony phrase h, where the new hypothesis is just h appended to hyp. All the new hypotheses generated in this iteration are added to S_{t+1}. After all the new hypotheses have been added, the hypotheses in S_{t+1} are sorted in descending order by their $S(hyp|m_0...m_{min(2t,l)})$ values, and only the top k are saved for the next iteration.

4.6 Generating Multiple Parts

The decoder was decribed above with the assumption that there is only one melody voice for which to create translations. In reality, when generating multiple harmony parts, the j^{th} harmony voice generated will be constrained by several source voices: the melody voice, M and the $j-1$ previous harmony voices generated. Let this set of source voices be defined as follows

$$S^j = M \cup \{ H^k | 1 \le k \le j-1 \}$$

In order to grow a harmony hypothesis in this setup, the melody phrase occurring after the end of the hypothesis to be grown must be determined for each source voice. From there, the harmony phrase possibilities can be retrieved and appended to the hypothesis as described above. Note that the 2-note phrases from the souce voices may be of varying durations, resulting in the hypotheses being of varying durations. As a consequence of this, it is possible that the end of a haromony hypothesis may not line up with the end of a note in one of the source voices. Therefore, the 2-note phrases that are translated may actually be versions of the notes in the source voices whose beginning note durations are altered so that the phrase begins only when the hypothesis ends. For example, if a hypothesis has duration 2.0 and, for some $s \in S^j$, the phrase that begins at or before 2.0 quarter-lengths, denote it (m_i, m_{i+1}), begins at 1.5 quarter-lengths, then the melody phrase to be translated from s will be (m_i, m_{i+1}), but (m_i) will have 0.5 quarter-lengths less of duration.

As mentioned earlier, the order in which the parts are generated is important since each part is contrained on all the parts generated before it. One ordering option is to choose some arbitrary ordering that is believed to be sufficient. However, there is no guarantee that an ordering that produces nice sounding harmonies for several compositions will produce optimal harmonies for all compositions. Instead of arbitrary choice, a greedy search for the best ordering is used to generate the parts.

In the greedy search, all harmony voices not yet generated are generated with the current source voices. The voice with

the highest weighted log-linear score is added to the source voices, and the process continues until all harmony voices to be generated are in the set of source voices. The intuition behind this method is that one would like to build up the musical arrangement with the strongest harmony voices possible. It stands to reason that constraining the next harmony voice with high-quality source voices will lead to a higher quality generation than one which is constrained on low-quality source voices.

5. SYSTEM PERFORMANCE

5.1 Model Evaluation Criteria

In order to evaluate the quality of the trained models, the perplexity metric, a common metric in natural language processing, is employed. This metric captures how well the algorithm scores arrangements in the test set. Perplexity PP [6] is defined over all compositions c in a test set as:

$$log_2 PP = -\sum_c log_2 S(H_c|M_c)$$

This measure provides a real number metric to evaluate different iterations of the system. Evaluations were done over 2 main test sets. One was a set of 28 major Bach chorales, pulled from the same corpus but not overlapping with the training data. The second was a set of 50 randomly generated compositions, referred to as anti-examples. The purpose of training on these two sets was to give evaluation scores similar to precision and recall. A high perplexity on the anti-examples test set implies good precision, whereas a low perplexity on the Bach test set implies good recall.

5.2 Major and Minor training data

Three systems trained on differing subsets of the training set were evaluated with the perplexity method. The first system used models trained entirely on chorales in the major mode, matching the test set of data. 185 different chorales were used for training. The second system used models trained on both major and minor chorales, and therefore had a larger data set of 351 songs to train on. The third systems used models trained entirely on minor chorales, of which there were 166 songs in our set, to act as a 'sanity check' for the evaluation algorithm.

**** PRECISION RECALL GRAPH FOR FULL COMPOSITION ****

Melody	Harmony	Major	Minor	Both
Soprano	Alto	27057.10	30821.11	25543.69
Soprano	Tenor	27929.47	32088.44	27359.35
Soprano	Bass	30157.27	34182.05	29438.73
Alto	Soprano	25746.65	28068.96	23830.84
Alto	Tenor	29965.21	32158.57	28588.83
Alto	Bass	33918.48	34980.79	31759.61
Tenor	Soprano	26590.96	29394.93	25802.08
Tenor	Alto	30489.94	31319.76	28608.72
Tenor	Bass	34761.60	35816.94	33013.67
Bass	Soprano	32062.73	34130.71	30866.27
Bass	Alto	37547.09	37652.04	35131.01
Bass	Tenor	38432.44	39039.53	36573.76

Table 3: The perplexity over the Bach test set of translation between each pair of voices for models trained on compositions in the major mode, the minor mode, and both modes.

Melody	Harmony	Major	Minor	Both
Soprano	Alto	27057.10	30821.11	25543.69
Soprano	Tenor	27929.47	32088.44	27359.35
Soprano	Bass	30157.27	34182.05	29438.73
Alto	Soprano	25746.65	28068.96	23830.84
Alto	Tenor	29965.21	32158.57	28588.83
Alto	Bass	33918.48	34980.79	31759.61
Tenor	Soprano	26590.96	29394.93	25802.08
Tenor	Alto	30489.94	31319.76	28608.72
Tenor	Bass	34761.60	35816.94	33013.67
Bass	Soprano	32062.73	34130.71	30866.27
Bass	Alto	37547.09	37652.04	35131.01
Bass	Tenor	38432.44	39039.53	36573.76

Table 4: The perplexity over the anti-examples test set of translation between each pair of voices for models trained on compositions in the major mode, the minor mode, and both modes.

other system in which all three models are used but are weighted equally, and lastly the final system described above. All three systems were trained on the same 185 major Bach chorales. Tables 5 and 6 show the perplexity results over the Bach and anti-example test sets, respectively.

****PRECISION RECALL GRAPH FOR FULL COMPOSITION****

5.3 Using Phrase Model and Tuned Weights

Another set of system iterations that were evaluated consisted of a system in which just the the note translation model and language model are used with equal weights, an-

5.4 Qualitative evaluation

Melody	Harmony	No Phrase Model	Phrase Model, Equal
Soprano	Alto	27057.10	30821.11
Soprano	Tenor	27929.47	32088.44
Soprano	Bass	30157.27	34182.05
Alto	Soprano	25746.65	28068.96
Alto	Tenor	29965.21	32158.57
Alto	Bass	33918.48	34980.79
Tenor	Soprano	26590.96	29394.93
Tenor	Alto	30489.94	31319.76
Tenor	Bass	34761.60	35816.94
Bass	Soprano	32062.73	34130.71
Bass	Alto	37547.09	37652.04
Bass	Tenor	38432.44	39039.53

Table 5: The perplexity over the Bach test set of translation between each pair of voices for three different iterations of the system.

Melody Equal Weights Tuned Weights	Harmony Phrase Model	No Phrase Model	Phrase M...	
Soprano	Alto	27057.10	30821.11	25943.09
Soprano	Tenor	27929.47	32088.44	27359.35
Soprano	Bass	30157.27	34182.05	29438.73
Alto	Soprano	25746.65	28068.96	23830.84
Alto	Tenor	29965.21	32158.57	28588.83
Alto	Bass	33918.48	34980.79	31750.61
Tenor	Soprano	26590.96	29394.93	25802.08
Tenor	Alto	30489.94	31319.76	28608.72
Tenor	Bass	34761.60	35816.94	33013.67
Bass	Soprano	32062.73	34130.71	30866.27
Bass	Alto	37547.09	37652.04	35131.04
Bass	Tenor	38432.44	39039.53	36573.76

Table 6: The perplexity over the anti-examples training set of translation between each pair of voices for three different iterations of the system.

Figure 3: The last 2 measures of a generated composition.

In Figure 1, we have the last two measures of one of the generated compositions from the most recent iteration of our algorithm. The Soprano part was given, and the rest of the voices were generated, with the Bass generated first, then Alto, then Tenor, as this seemed subjectively to produce the best result. The major language and translation models were used for the same reason. Each note will be referred to by the first letter of its voice, followed by the index of the note in the voice (e.g. S3, T5). As is conventional, the Soprano and Alto parts are written in treble clef, while the Tenor and Bass parts are written in bass clef.

5.4.1 Motion

In this example, we have several instances of contrary motion. Contrary motion occurs when the notes in two voices move in opposite directions, and is generally regarded as ideal in counterpoint. In our example, we can see this occuring in S1-S2 and T1-T2, in A2-A3 and B2-B3, and in several other places.

We can also see examples of oblique motion, in which one voice moves while the other stays the same. This is regarded as good, although not as good as contrary motion. We see examples of this in S2-S3 and A2-A3, in A4-A5 and T4-T5, and in several other places.

Finally, we have some examples of similar motion, which is when notes move in the same direction, and parallel motion, which is similar motion where the notes move exactly the same distance. Parallel motion is considered the worst kind of motion, with similar motion being slightly better. Some amount of similar motion cannot be avoided in compositions with more than three parts. However, we have examples of parallel motion betwween S3-S4 and A3-A4 and T3-T4, which is not ideal. This suggests that improvements could be made to our algorithm to pay closer attention to motion in order to avoid these outputs.

5.4.2 Harmonies

The chords comprised of each of the voices in this example actually make the chords C (1), G (2), Dm7 (3), G (4), and C (5). This is interesting because there was no part of the algorithm deciding what chord to generate at each step, only the individual notes in relation to each other. Regardless, it created this set of chords, which is a reasonable chord progression to find in a piece of music similar to the data we trained on. Additionally, the last three chords comprise a ii-V-I (two-five-one) cadence, which is a very common way to end compositions.

6. ETHICS

As with most artificial intelligence applications, an obvious ethical concern with this system is its potential impact on people whose activities it attempts to mimic. If MelTS produced compositions with such high quality that they were competitive with the compositions of modern professional composers, the system may threaten the livelihoods of composers and musicians. However, the authors believe that

such a high level of output quality is so far in the future that this hypothetical situation could not be considered a legitimate concern.

If MelTS were to be widely used as an aid to musicians, there would be some foreseeable legal issues regarding ownership of the compositions. For example, would the composers of the training data have a right to the output? Would the output belong to the composer of the input harmony? To the creators of MelTS? Since music is generally considered intellectual property, it is clear that ownership of the output would need to be delineated to all users and composers of training data, but what the delineation might be is not immediately obvious.

7. REMAINING WORK

A natural next step for MelTS would be to release the system to the public, preferably as a web application that would allow users to generate harmonies based on pre-computed models. Musicians could use the tool to experiment with harmonization ideas by exploring the way different genres of music would harmonize their melodies. Additionally, feedback from real users of the application would allow for further improvements to the system.

In order to publish a viable a web app, the generation algorithm would need to be optimized for time efficiency. Minimal work has been done to systematically identify bottlenecks in the algorithm, although parallelization appears to be a clear way of speeding up. For example, generating all remaining parts based on the previously generated parts and extending harmony hypotheses during decoding can all be done in parallel.

Another future goal is to train more models in different styles of music. Potential styles would be marching band arrangements, orchestral scores, and gospel music. With some creativity in building the models, MelTS could even be used to generate chord accompaniment for pop songs.

As for improvements to the system itself, it would be interesting to investigate adding more features to the weighted log-linear score. For example, the chords for a beat could be predicted and incorporated into the score through a feature that captures whether or not the harmonized notes occur in one of the top predicted chords. To maintain a legal range of pitches within parts, a feature representing a note's distance from the center of the pitch range for a given part could also be incorporated.

8. REFERENCES

[1] David Cerny, Jiten Suthar, and Israel Geselowitz. U-amp: User input based algorithmic music platform. 2014.

[2] Michael Scott Cuthbert and Christopher Ariza. music21: A toolkit for computer-aided musicology and symbolic music data.

[3] Kemal Ebcioğlu. An expert system for harmonizing chorales in the style of j.s. bach. *The Journal of Logic Programming*, 8(1 - 2):145 – 185, 1990. Special Issue: Logic Programming Applications.

[4] Hermann Hild, Johannes Feulner, and Wolfram Menzel. Harmonet: A neural net for harmonizing chorales in the style of j. s. bach. In J.E. Moody, S.J. Hanson, and R.P. Lippmann, editors, *Advances in Neural Information Processing Systems 4*, pages 267–274. Morgan-Kaufmann, 1992.

[5] L.A.J. Hiller and L.M. Isaacson. *Experimental Music: Composition with an Electronic Computer*. McGraw-Hill, 1959.

[6] Philipp Koehn. *Statistical Machine Translation*. Cambridge University Press, New York, NY, USA, 1st edition, 2010.

[7] Franz Josef Och and Hermann Ney. Statistical multi-source translation. *MT Summit 2001*, pages 253–258, 2001.

[8] Stanisław A. Raczyński, Satoru Fukayama, and Emmanuel Vincent. Melody harmonization with interpolated probabilistic models. *Journal of New Music Research*, 42(3):223–235, 2013.

[9] Jordi Sabater, J. L. Arcos, and R. López De Mántaras. Using rules to support case-based reasoning for harmonizing melodies. In *Multimodal Reasoning: Papers from the 1998 AAAI Spring Symposium (Technical Report, WS-98-04). Menlo Park, CA*, pages 147–151. AAAI Press, 1998.

[10] R. Scholz, E. Vincent, and F. Bimbot. Robust modeling of musical chord sequences using probabilistic n-grams. In *Acoustics, Speech and Signal Processing, 2009. ICASSP 2009. IEEE International Conference on*, pages 53–56, April 2009.

[11] Ian Simon, Dan Morris, and Sumit Basu. Mysong: Automatic accompaniment generation for vocal melodies. In *CHI 2008 Conference on Human Factors in Computing Systems*. Association for Computing Machinery, Inc., April 2008.

[12] C.P. Tsang and M. Aitken. Harmonizing music as a discipline of constraint logic programming. 1991.